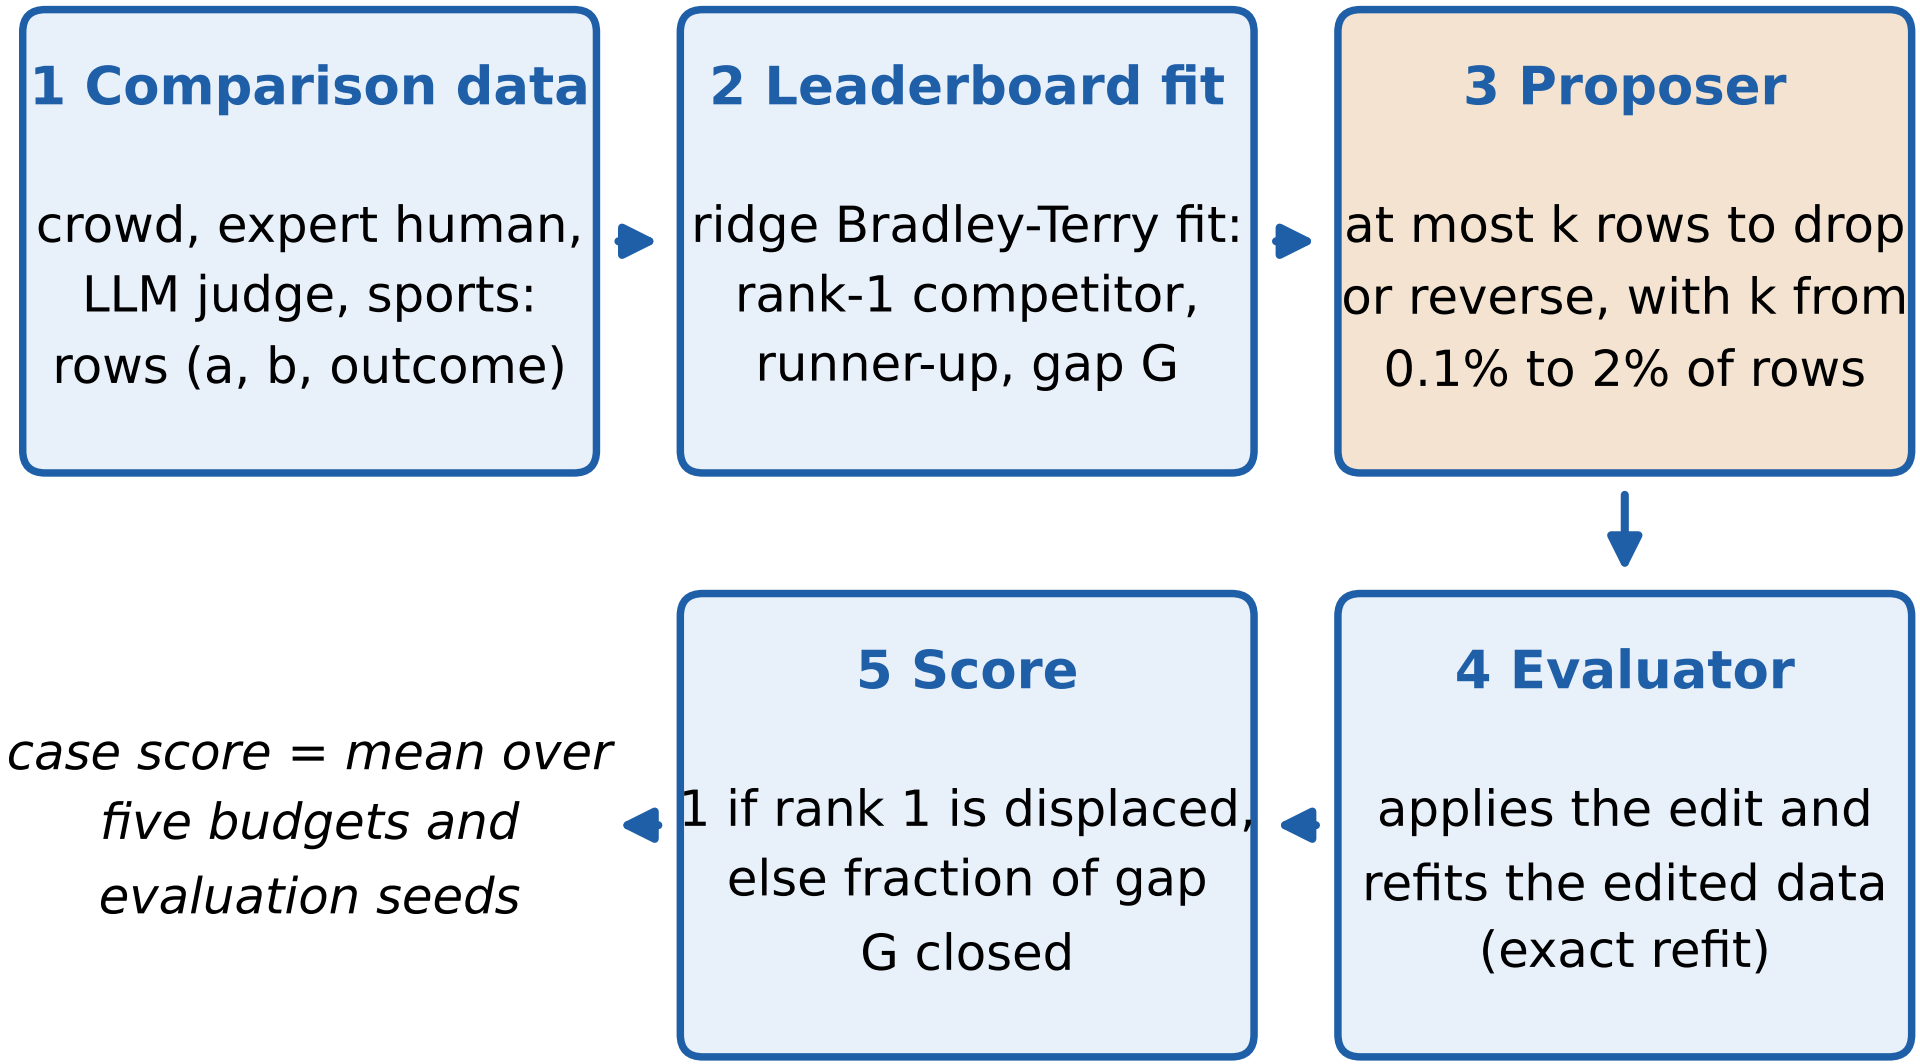
1 Comparison data
crowd, expert human, LLM judge, sports: rows (a, b, outcome)
2 Leaderboard fit
ridge Bradley-Terry fit: rank-1 competitor, runner-up, gap G
3 Proposer
at most k rows to drop or reverse, with k from 0.1% to 2% of rows
4 Evaluator
applies the edit and refits the edited data (exact refit)
5 Score
1 if rank 1 is displaced, else fraction of gap G closed
case score = mean over five budgets and evaluation seeds