Table 1: Measured completed-job costs, separate from the initial pilot and resource projections. CPU and wall columns give the successful completion process time; reference CPU sums all saved exact-reference call timings. These scopes overlap and must not be added. A resumed process can exclude earlier work, which belongs in the full research compute ledger alongside pilots and failed attempts. Peak memory is the completion-process peak RSS. The independent saved-panel audit records 0 failed attempts, 0 warning messages and 0 exhaustive calls among the nominally stochastic runs; those counts must accompany any final efficiency or stochastic-method claim. Severe execution interruptions contaminate wall times; they are preserved execution records, not deployment-latency benchmarks. Use CPU and actual rows for efficiency comparisons. Runtime excludes untracked organizational work and is not a per-call latency quantile.

Dataset	Model	CPU min	Wall min	Ref. CPU min	Peak MiB
German	LR	24.63	29.32	17.589	499
German	RF	175.80	2905.87	151.285	521
German	HGB	195.78	2926.88	169.846	486
ACSIncome	LR	19.04	19.07	0.299	812
ACSIncome	RF	16.83	21.55	0.285	436
ACSIncome	HGB	8.06	8.07	0.158	436
Diabetes	LR	0.78	0.79	0.001	314
Diabetes	RF	5.35	5.36	0.035	314
Diabetes	HGB	3.33	3.35	0.045	314