Table 1: Measured pilot CPU time per explanation call (milliseconds, three significant figures). C1 and C2 are the two pilot inputs per fitted model; each number is one timed call, not an average or confidence interval. Column labels are requested coalition-work settings, not actual prediction-row counts. The background contains 32 rows. Library import was warmed before timing; training and references are excluded. The initial 2048 Kernel setting is exhaustive for ACSIncome and Diabetes and is retained as a labeled control, not stochastic main-study evidence. LR, RF and HGB denote logistic regression, random forest and histogram gradient boosting. This table describes the pilot, not final-study costs.

			Permutation (ms)		Kernel (ms)	
Dataset	Model	Input	128	2048	128	2048
German	LR	C1	2.24	35.0	15.4	163
		C2	1.57	23.8	10.6	132
	RF	C1	24.2	291	47.7	280
		C2	22.9	291	47.1	290
	HGB	C1	19.7	307	32.6	362
		C2	20.9	299	29.1	344
ACSIncome	LR	C1	31.1	223	67.8	516
		C2	29.9	203	73.1	532
	RF	C1	46.4	235	78.5	284
		C2	40.4	253	79.7	295
	HGB	C1	16.4	106	27.8	154
		C2	16.8	129	26.1	137
Diabetes	LR	C1	1.13	4.90	5.50	5.28
		C2	0.930	4.46	5.08	5.28
	RF	C1	20.5	39.3	40.1	51.1
		C2	19.7	41.0	46.6	59.6
	HGB	C1	17.6	51.8	31.7	53.4
		C2	15.2	42.8	29.6	46.9